

# Memcpy to Meaning: a 64 kB SRAM Slab that Routes Tools in Real Time without Tokens, Back Prop, or Regret

## 1. Introduction

The rapid proliferation of generative artificial intelligence, particularly large language models (LLMs), has fundamentally reshaped the landscape of human-computer interaction. These models, capable of producing coherent text, images, and code, have unlocked unprecedented capabilities in content creation, problem-solving, and conversational agents. However, this power comes at a significant cost. The computational, energetic, and latency overheads associated with state-of-the-art generative models present substantial barriers to their deployment in real-time, resource-constrained, or interactive applications. The very architecture that enables their generative prowess—the autoregressive, token-by-token generation process—introduces inherent inefficiencies that are becoming increasingly difficult to overcome through hardware acceleration alone. This paper introduces Neural Pin, a novel system designed to operate in a "post-generative" paradigm, where the goal is not to generate new content from scratch but to efficiently route and orchestrate existing tools based on a learned, stateful understanding of user intent. By fundamentally rethinking the nature of inference, Neural Pin sidesteps the bottlenecks of token-based systems, offering a path to real-time, energy-efficient, and highly responsive AI agents.

### 1.1 The Bottlenecks of Generative AI

The current generation of AI systems, dominated by transformer-based architectures, faces several critical bottlenecks that limit their applicability in dynamic, interactive environments. These challenges span computational cost, memory management, and energy consumption, creating a gap between the potential of AI and its practical, real-world implementation.

#### 1.1.1 The Computational Cost of Token-Based Inference

The core mechanism of generative models like GPT-4 is autoregressive token generation. Each new token is predicted based on the entire preceding sequence, a process that requires a full forward pass through a massive neural network. This sequential dependency means that generation cannot be parallelized, creating a fundamental latency floor for each output token. While techniques like batching can improve throughput for multiple concurrent requests, they do not reduce the latency for any single request, which is critical for real-time interaction. Furthermore, the computational complexity of the attention mechanism within the transformer architecture scales quadratically with the sequence length, making long-context interactions prohibitively expensive. This token-by-token approach, while powerful for generating novel content, is inherently inefficient for tasks that require rapid, stateful decision-making rather than open-ended creation. The energy cost of this process is also substantial; a single query to a model like ChatGPT is estimated to consume approximately five times more electricity than a standard web search, a figure that is expected to grow as models become larger and more complex. This high computational and energetic cost makes it impractical to run sophisticated generative models locally on commodity hardware or in environments with strict power budgets.

#### 1.1.2 Memory and Latency Challenges in Transformer Models

A significant challenge in deploying large language models is managing the immense memory footprint required for inference. The model parameters themselves can occupy tens or even hundreds of gigabytes, but an even more pressing issue is the memory required for the key-value (KV) cache. During the prefill stage of inference, where the initial prompt is processed, the model computes and stores key-value pairs for every token in the context. This KV cache must be retained in high-speed memory (typically GPU VRAM) for the entire decoding stage, as it is repeatedly accessed to generate each subsequent token. The size of the KV cache grows linearly with both the batch size and the sequence length, quickly exhausting the available memory of even high-end accelerators. For example, for a batch size of 512 and a context length of 2048, the KV cache can consume up to 3 terabytes of memory, which is three times the size of the model itself. This "memory wall" is a primary bottleneck that limits the maximum context length and the number of concurrent requests a system can handle. While recent work, such as the "Pie" system, has explored techniques like performance-transparent swapping to extend GPU memory using high-bandwidth CPU memory, these solutions add architectural complexity and can introduce latency if not managed perfectly. The need to constantly manage and access this large, growing cache contributes significantly to the overall latency and energy consumption of the inference process, making it difficult to achieve the near-instantaneous response times required for fluid human-computer interaction.

#### 1.1.3 The Overhead of KV Caches in Large Language Models

The Key-Value (KV) cache is a critical optimization for autoregressive transformer inference, but it also represents a major source of overhead. Its purpose is to store the intermediate attention states (keys and values) for all tokens in the context, preventing the need to recompute them for every new token generated. However, this benefit comes at the cost of a large and dynamically growing memory requirement. For a model with billions of parameters and a long context window, the KV cache can easily consume tens of gigabytes of GPU memory. This not only increases the hardware cost but also limits the scalability of the system. The management of the KV cache is a complex task for inference engines like vLLM, TGI, and TensorRT-LLM, which must orchestrate its allocation and access efficiently across multiple concurrent requests. Furthermore, the size of the cache directly impacts the latency of each decoding step, as a larger cache requires more memory bandwidth to access. The "Pie" system addresses this by using high-bandwidth interconnects like NVLink to swap portions of the cache between CPU and GPU memory, but this introduces a new layer of complexity and a potential performance bottleneck if the prefetching of data is not perfectly synchronized with the computation. This overhead is a direct consequence of the token-based architecture, where the entire history of the conversation must be explicitly stored and accessed for every new decision.

#### 1.2 A New Paradigm: Post-Generative Computing

In response to the limitations of generative AI, we propose a shift towards a "post-generative" computing paradigm. This approach does not seek to replace generative models but to complement them by providing a more efficient mechanism for tasks that do not require open-ended content creation. The core idea is to move beyond the

token-generation framework and treat inference as a stateful, continuous process of updating a compact representation of intent.

### 1.2.1 Moving Beyond Token Generation

The post-generative paradigm is predicated on the observation that many intelligent behaviors, particularly in interactive systems, can be modeled as a continuous process of state updates rather than discrete acts of generation. Instead of generating a long sequence of tokens to describe a plan or an intent, the system maintains a rich, high-dimensional vector that represents the current state of its "mind." This vector is not a static embedding but a dynamic state that is continuously updated as new information arrives. When a tool produces a result, that result is not parsed into tokens and fed into a prompt; instead, its semantic content is directly integrated into the state vector through a simple, efficient update rule. This approach eliminates the need for the complex and costly autoregressive generation process for every decision. The system's "understanding" is not stored in a sequence of words but in the configuration of this compact vector state. This allows for a much more fluid and responsive interaction model, where the system's intent can evolve in real-time without the overhead of re-generating a context window.

### 1.2.2 Inference as a Stateful Memory Write

In the post-generative model, the act of inference is re-conceptualized as a stateful memory write. The "Mind" of the system is a fixed-size block of memory (analogous to an SRAM bank) that holds the current state vector. When new information arrives, it is not processed through a multi-layer transformer but is instead used to perform a simple, in-place update on this memory. This update, implemented using SIMD instructions for maximum efficiency, is a form of online learning where the state vector is nudged in the direction of the new information. This process is extremely fast and energy-efficient, as it involves only a small number of floating-point operations per dimension. The "inference" is not a query that returns a result but an action that modifies the system's internal state. This stateful approach means that the system's understanding is cumulative and persistent within a session, without the need to maintain a growing history of tokens. The entire "context" is compressed into this fixed-size vector, making it highly memory-efficient and eliminating the need for complex cache management.

### 1.2.3 The Vision of a Token-Free, Gradient-Free System

The ultimate vision of the post-generative paradigm is a system that is both token-free and gradient-free. "Token-free" means that the internal representation of information is not based on a discrete vocabulary of tokens but on continuous, high-dimensional vectors. This eliminates the computational overhead of tokenization, detokenization, and the sequential processing of token sequences. "Gradient-free" refers to the learning mechanism. Instead of using backpropagation to compute gradients across a deep network, the system uses a simple, local update rule to modify its state vector. This is a form of Hebbian or contrastive learning, where the state is updated based on the similarity between the current state and the new input. This approach is vastly more efficient than gradient-based optimization, as it does not require storing intermediate activations or performing a backward pass. The combination of these two properties

results in a system that is incredibly lightweight, fast, and energy-efficient, making it suitable for deployment on a wide range of devices, from powerful servers to resource-constrained edge devices.

### 1.3 Contributions of Neural Pin

Neural Pin is a concrete implementation of the post-generative computing paradigm. It is a system designed to demonstrate the feasibility and advantages of this new approach in a practical application: real-time tool routing. The contributions of Neural Pin are threefold, spanning system design, implementation, and performance.

#### 1.3.1 A Lock-Free, Token-Less Semantic Router

The primary contribution of Neural Pin is its design as a lock-free, token-less semantic router. It provides a mechanism for an AI agent to learn a user's intent and select the appropriate tool to execute next, without ever generating or processing a single token. The system's "Mind" is a 64 kB `SharedArrayBuffer` that holds a 512-dimensional vector representing the current state. Tool results, also represented as 512-dimensional vectors, are streamed into a circular ring buffer. A dedicated "Spindle" component, running asynchronously, rotates the ring buffer and triggers an update to the "Mind." This entire process is lock-free, relying on atomic operations to ensure thread safety, which eliminates the performance bottlenecks and complexity associated with traditional locking mechanisms. The semantic routing is achieved by monitoring the cosine similarity between the "Mind" vector and a set of predefined "Facets," which represent different potential intents. When the similarity to a facet crosses a learned threshold, a "Pin" is triggered, and the system selects the corresponding tool.

#### 1.3.2 A 64 kB Shared Memory System for Real-Time Tool Routing

Neural Pin demonstrates that a surprisingly small amount of memory is sufficient for sophisticated, real-time decision-making. The entire system, including the "Mind," the ring buffer, and all necessary metadata, fits within a single 64 kB `SharedArrayBuffer`. This fixed memory footprint is a stark contrast to the gigabytes of memory required by generative models. The use of a `SharedArrayBuffer` allows for efficient, zero-copy communication between the main application thread and the worker thread that manages the "Mind." This shared memory approach, combined with the lock-free design, enables extremely low-latency updates and decision-making. The system is designed to be multi-tenant, with each user or session having its own isolated 64 kB "Mind" slab, all sharing a common ring buffer for tool results. This architecture makes it highly scalable and suitable for deployment in multi-user environments.

#### 1.3.3 A Post-Generative Model with Sub-Millisecond Convergence

The performance of Neural Pin validates the post-generative approach. The system is designed to converge on a user's intent in a remarkably short amount of time. The abstract convergence in "12 human-time clicks," which is approximately 120 milliseconds on a commodity laptop. This is achieved through a highly optimized update loop. The in-place gradient step, which updates the "Mind" vector, is implemented using SIMD instructions and is designed to take less than a microsecond. The energy cost per routing decision is claimed to be less than 1 millijoule, a figure that is orders of magnitude lower than the energy consumption of generative models. The entire system is implemented in just 280 lines of TypeScript and WebAssembly,

demonstrating the elegance and simplicity of the post-generative paradigm. These performance characteristics make Neural Pin a compelling alternative for applications where real-time responsiveness and energy efficiency are paramount.

## 2. System Architecture

The architecture of Neural Pin is a direct reflection of its post-generative philosophy: simple, efficient, and stateful. It is composed of a small number of core components that work together to create a system that is both powerful and lightweight. The design draws inspiration from hardware concepts like SRAM, DMA controllers, and CPU interrupts, translating them into a software implementation that is optimized for modern multi-core processors.

### 2.1 Core Components

The system is built around four main components, each with a specific role in the data flow and decision-making process. These components are designed to be highly cohesive and loosely coupled, communicating through a shared memory space to achieve lock-free operation.

#### 2.1.1 The Mind: A 64 kB Mutable Vector State

Name Hardware analogue Purpose

Mind SRAM bank 64 kB mutable vector that is the current thought

Spindle DMA controller rotates the ring (async, lock-free)

Pin CPU interrupt wakes the Mind when a facet crosses threshold

Facet register file 8–16 named 512-D sub-vectors inside the Mind

The "Mind" is the heart of the Neural Pin system. It is a 64 kB block of memory, implemented as a `SharedArrayBuffer`, that holds the current state of the system's "thought." This state is represented as a single, high-dimensional vector of 512 floating-point numbers. This vector is not a static representation but a dynamic one, constantly evolving as the system receives new information. The "Mind" is analogous to an SRAM bank in a hardware system, providing fast, in-place read and write access to the current state. The entire learning and inference process of Neural Pin revolves around updating this single vector. The 64 kB size is a deliberate choice, providing enough space for a rich 512-dimensional representation while remaining small enough to fit in the CPU's L1 or L2 cache, ensuring extremely fast access times. The "Mind" is a purely mutable state; there are no immutable copies or versions, reflecting the continuous, flowing nature of the post-generative paradigm.

#### 2.1.2 The Spindle: A Lock-Free Ring Buffer Controller

The "Spindle" is the component responsible for managing the flow of data into the system. It acts as a DMA (Direct Memory Access) controller, asynchronously moving data from an input source into the "Mind." The input source is a circular ring buffer, also implemented within the `SharedArrayBuffer`, that holds the results from various tools. Each tool result is a 512-float vector. The Spindle's job is to "rotate" this ring buffer, which involves atomically updating a head pointer to point to the next available slot. This operation is performed using a lock-free atomic compare-and-swap (CAS) instruction, ensuring that multiple producers (tool-executing threads) can write to the ring buffer without contention. When the Spindle rotates the buffer, it also triggers the update process for the "Mind," taking the vector from the newly consumed slot and

using it to modify the "Mind's" state. The asynchronous, non-blocking nature of the Spindle is crucial for maintaining the real-time performance of the system.

### 2.1.3 The Pin: A Cosine-Triggered Interrupt

The "Pin" is the system's decision-making trigger. It is analogous to a CPU interrupt that fires when a specific condition is met. In Neural Pin, this condition is a semantic one: the cosine similarity between the "Mind" vector and one of the predefined "Facets" crosses a certain threshold. The system continuously monitors the "Mind" vector, calculating its similarity to a small set of "Facets," which are essentially named sub-vectors within the "Mind" that represent different potential user intents (e.g., "search for information," "summarize a document," "write code"). When the similarity to a particular facet becomes high enough, the "Pin" fires, signaling that the system has recognized a strong intent. This trigger is not just a binary signal; it also implicitly updates the "intent" facet of the "Mind," preparing it for the next step. The threshold for triggering the Pin is adaptive, allowing the system to learn and adjust its sensitivity over time. This cosine-triggered interrupt mechanism is the core of the system's semantic routing capability, allowing it to make decisions based on the high-level state of the "Mind" rather than on explicit, token-based commands.

### 2.1.4 Facets: Named Sub-Vectors for Semantic Routing

"Facets" are a key abstraction within the "Mind." They are not separate vectors but are defined as fixed offsets within the single 512-dimensional "Mind" vector. For example, the first 64 dimensions might be designated as the "search" facet, the next 64 as the "summarize" facet, and so on. These facets represent different semantic categories or potential user intents. The system learns to associate certain patterns in the input data with these facets. The "Pin" mechanism works by calculating the cosine similarity between the entire "Mind" vector and the vector formed by each facet's dimensions. This allows the system to have a multi-faceted understanding of the user's intent, where multiple facets can be active to varying degrees at the same time. The "intent" facet is a special facet that is updated when a Pin fires, effectively "locking in" the most likely intent and preparing the downstream router to select the appropriate tool. This design allows for a rich, nuanced representation of intent within a compact, fixed-size vector.

## 2.2 Physical Layout and Multi-Tenancy

The physical layout of Neural Pin is designed for efficiency and scalability. It leverages modern web technologies to create a system that is both powerful and portable.

### 2.2.1 SharedArrayBuffer as the Foundation

The entire Neural Pin system is built on top of a single SharedArrayBuffer. This is a JavaScript object that represents a generic, fixed-length raw binary data buffer that can be shared between multiple threads. In the context of Neural Pin, the SharedArrayBuffer is the 64 kB "slab" of memory that contains the "Mind," the ring buffer, and all the necessary control structures. Using a SharedArrayBuffer provides several key advantages. First, it allows for zero-copy communication between the main application thread and the worker thread that runs the "Mind" update loop. The worker can directly read and write to the same memory space as the main thread, eliminating the need for costly data serialization and deserialization. Second, it provides the


foundation for the lock-free design. The Atomics API in JavaScript provides atomic operations like `compareExchange` that can be used to safely coordinate access to the shared memory without traditional locks. This is essential for achieving the high performance and low latency that are core goals of the system.

### 2.2.2 A Circular Ring Buffer for Streaming Tool Results

The input to the Neural Pin system is a stream of tool results, each represented as a 512-float vector. These results are fed into a circular ring buffer that is located within the `SharedArrayBuffer`. The ring buffer has a fixed number of slots (e.g., 32), and it operates on an overwrite-only principle. When a new tool result arrives, it is copied into the next available slot in the ring buffer using a `memcpy`-like operation. The "Spindle" then atomically updates the head pointer to point to this new slot, effectively consuming the data. If the buffer is full, the oldest data is simply overwritten. This design has several benefits. It is extremely simple and efficient, with no need for complex memory management or garbage collection. The overwrite-only nature ensures that the system can handle a high throughput of tool results without ever blocking or running out of memory. It also provides a natural form of temporal decay, where older results are gradually forgotten as they are overwritten by newer ones.

### 2.2.3 Multi-Tenant Support with Per-Session Minds

Neural Pin is designed to be a multi-tenant system, capable of supporting multiple users or sessions simultaneously. This is achieved by creating a separate "Mind" for each tenant. Each "Mind" is its own isolated 64 kB `SharedArrayBuffer`, containing its own state vector, ring buffer, and control structures. These individual "Mind" slabs can all share a common, global ring buffer for incoming tool results, or each can have its own dedicated input stream. The key is that the state of each user's interaction is completely isolated from the others. This per-session design is highly scalable, as new users can be accommodated simply by allocating a new 64 kB block of memory. The lightweight nature of each "Mind" means that a single server can potentially support thousands of concurrent sessions without a significant memory footprint. This makes Neural Pin an ideal architecture for building large-scale, interactive AI services.



▲ ▲ ▲

## 2.3 Operational Flow

The operation of Neural Pin is a continuous, cyclical process that can be broken down into a series of steps. This flow is designed to be as efficient as possible, with each step optimized for speed and minimal resource consumption.

### 2.3.1 Step 1: Tool Result Ingestion via `memcpy`

The process begins when a tool, such as a web search or a code executor, completes its task and produces a result. This result is converted into a 512-dimensional floating-point vector, which captures its semantic meaning. This vector is then copied directly into the next available slot in the circular ring buffer. This operation is a simple, fast memory copy ( `memcpy` ), with no parsing, tokenization, or other processing required.

The address of the slot is determined by the current head pointer of the ring buffer, which is managed by the "Spindle." This step is designed to be as lightweight as possible, allowing the tool-executing thread to quickly return to its task.

#### 2.3.2 Step 2: Asynchronous Ring Rotation by the Spindle

Once the tool result has been copied into the ring buffer, the "Spindle" takes over. Running in a separate worker thread, the Spindle's job is to "rotate" the ring buffer. This is done by atomically incrementing the head pointer using a compare-and-swap (CAS) operation. This ensures that even if multiple tool threads are writing to the buffer simultaneously, the head pointer is updated in a consistent and thread-safe manner. The rotation of the ring buffer signals that a new piece of data is ready to be processed. The Spindle then takes the vector from the slot that was just consumed and uses it to trigger the update of the "Mind."

#### 2.3.3 Step 3: Pin Firing on Cosine Similarity Threshold

As the "Mind" vector is updated with the new information, the "Pin" mechanism is constantly monitoring its state. It calculates the cosine similarity between the current "Mind" vector and each of the predefined "Facets." If the similarity to any facet crosses its adaptive threshold, the "Pin" fires. This is the moment of "insight" for the system,  
| Spindle | Spindle | Spindle



Shared 64 kB Ring Buffer (mmap) |

where it recognizes a strong, learned pattern in the user's intent. The firing of the Pin is an interrupt-like event that signals the downstream router that a decision has been made. The threshold itself is not static; it is learned over time, allowing the system to become more or less sensitive to certain intents based on the user's feedback and behavior.

#### 2.3.4 Step 4: In-Place Gradient Update of the Mind

When the Spindle consumes a new tool result vector, it triggers an in-place update of the "Mind" vector. This is not a traditional backpropagation step but a simple, efficient online learning rule. The update is performed using SIMD (Single Instruction, Multiple Data) instructions to maximize computational throughput. The update rule is of the form  $\text{target} += \text{lr} \cdot \delta + \text{momentum}$ , where target is the "Mind" vector, lr is the learning rate,  $\delta$  is the difference between the new tool result vector and the current "Mind" vector, and momentum is a term that helps to smooth the updates and prevent oscillations. This entire operation is designed to be extremely fast, taking less than a microsecond to complete. It is this rapid, in-place update that allows the "Mind" to continuously and efficiently absorb new information.

#### 2.3.5 Step 5: Implicit Intent Update and Tool Selection

The final step in the cycle is the selection of the next tool. When a "Pin" fires, it not only signals a decision but also implicitly updates the "intent" facet of the "Mind." This facet is then read by a downstream "RouterAgent." The RouterAgent does not need to parse any text or understand a complex prompt; it simply reads the 512-float vector of the "intent" facet and uses it to select the next tool to execute. This selection could be



based on a simple lookup table, a small neural network, or another lightweight routing mechanism. Once the tool is selected and executed, its result is fed back into the ring buffer, and the cycle begins anew. This closed-loop system allows for a continuous, adaptive, and highly efficient interaction between the user and the AI agent.

### 3. Methodology

The methodology of Neural Pin is centered on achieving maximum efficiency and real-time performance through a combination of low-level programming techniques, lock-free data structures, and a simplified learning model. The design philosophy is to do the minimum amount of work necessary to achieve the goal of semantic routing, leveraging hardware features like SIMD and atomic operations to accelerate every step of the process.

#### 3.1 The Mind: In-Place Vector Updates

The core of the Neural Pin system is the "Mind," a 64 kB SharedArrayBuffer that holds the system's state as a 512-dimensional vector. The learning process is not a complex backpropagation through a deep network but a simple, efficient, in-place update of this vector. This approach is designed to be extremely fast and to have a fixed, predictable memory footprint.

##### 3.1.1 SIMD Operations for Efficient Computation

To maximize the speed of the vector updates, Neural Pin leverages Single Instruction, Multiple Data (SIMD) operations. SIMD allows a single CPU instruction to operate on multiple data points simultaneously, which is ideal for the element-wise arithmetic required for vector addition and scaling. The update rule  $\text{target} += \text{lr} \cdot \delta + \text{momentum}$  involves several vector operations that can be significantly accelerated with SIMD. By processing multiple vector elements in parallel, the system can reduce the time required for a single update to well under a microsecond. This is a critical optimization that enables the real-time performance of the system, allowing the "Mind" to absorb new information almost instantaneously.

##### 3.1.2 The Update Rule: $\text{target} += \text{lr} \cdot \delta + \text{momentum}$

The learning rule used to update the "Mind" vector is a variation of the classic gradient descent formula with momentum. The formula is  $\text{target} += \text{lr} \cdot \delta + \text{momentum}$ , where: target is the 512-dimensional "Mind" vector being updated.

lr is the learning rate, a small constant that controls the size of the update step.

$\delta$  is the error vector, calculated as the difference between the new tool result vector and the current "Mind" vector.

momentum is a term that incorporates a fraction of the previous update, which helps to smooth out the learning process and accelerate convergence in a consistent direction.

This simple, local update rule is a form of online learning that allows the system to continuously adapt to new data without the need for a complex, multi-layer network or a global loss function.

##### 3.1.3 Zero Heap Allocations and Fixed Memory Footprint

A key design principle of Neural Pin is its zero-allocation policy. After the initial 64 kB SharedArrayBuffer is allocated during the system's cold start, no further dynamic memory allocations are performed on the heap. This has several important benefits.

First, it eliminates the overhead and potential latency spikes caused by garbage collection, which is a common issue in managed runtimes like JavaScript. Second, it ensures that the system's memory usage is constant and predictable, a critical requirement for real-time and embedded systems. The fixed memory footprint of 64 kB makes the system highly portable and easy to deploy in a wide range of environments.

### 3.2 The Spindle: A Lock-Free Ring Buffer

The Spindle is the component responsible for managing the flow of data into the system. It operates on a circular ring buffer, a classic data structure for high-throughput, real-time systems. The key innovation of the Spindle is its lock-free implementation, which allows for asynchronous, non-blocking operation.

#### 3.2.1 Implementation with Atomic Compare-and-Swap (CAS)

The Spindle uses an atomic compare-and-swap (CAS) operation to manage the head pointer of the ring buffer. The CAS operation is a fundamental building block of lock-free programming. It works by checking if the value of a memory location is equal to an expected value; if it is, the operation updates the memory location to a new value. This entire operation is atomic, meaning it cannot be interrupted by other threads. This allows multiple producer threads to safely update the head pointer without the need for a mutex, which would introduce blocking and potential deadlocks.

#### 3.2.2 Overwrite-Only Circular Buffer Design

The ring buffer is designed to be overwrite-only. It has a fixed number of slots (e.g., 32), and when the head pointer reaches the end, it wraps around to the beginning, overwriting the oldest data. This design is simple and efficient, as it avoids the need for complex memory management to free up space. It also prioritizes recency, ensuring that the system is always working with the most up-to-date information. This is appropriate for a real-time system where the latest data is often the most relevant.

#### 3.2.3 Asynchronous and Non-Blocking Operation

The Spindle operates asynchronously in a separate worker thread. This decouples the process of data ingestion from the main processing loop of the "Mind." When a tool result is written to the ring buffer, the producer thread can immediately return to its task without waiting for the "Mind" to process the data. This non-blocking operation is crucial for maintaining high throughput and low latency, as it allows the system to handle a continuous stream of incoming data without being bottlenecked by the processing speed of the "Mind."

### 3.3 The Pin: A Cosine-Triggered Interrupt

The Pin is the event-driven component of the system, responsible for triggering the "Mind" to process new information. It acts as a semantic filter, ensuring that the "Mind" only updates its state when it encounters data that is relevant to its current goals.

#### 3.3.1 Cosine Similarity as a Semantic Trigger

The Pin uses cosine similarity to measure the semantic relatedness between the incoming tool result vector and the "Facets" within the "Mind." Cosine similarity is a measure of the angle between two vectors, and it is a standard metric for comparing high-dimensional embeddings. A high cosine similarity score indicates that the two vectors are pointing in a similar direction in the semantic space, suggesting that the new information is relevant to a particular facet of the system's understanding.

### 3.3.2 Adaptive Thresholding for Dynamic Learning

The threshold for triggering the Pin is not a fixed value but is adaptive. It is a learned parameter that can be adjusted over time based on the system's experience. This allows the system to become more or less sensitive to different types of information. For example, if the system is consistently receiving irrelevant data, it can raise the threshold to become more selective. Conversely, if it is missing important information, it can lower the threshold to become more sensitive. This dynamic learning capability allows the system to adapt to different contexts and user behaviors.

### 3.3.3 Sub-Microsecond Wake-Up Time for the Mind

When the Pin fires, it wakes up the "Mind" to perform its update. This wake-up process is designed to be extremely fast, with a latency of less than a microsecond. This is achieved by using a simple, efficient signaling mechanism, such as an atomic flag or a promise, to notify the "Mind" worker thread that new data is ready for processing. This sub-microsecond wake-up time is a critical component of the system's real-time performance, ensuring that it can react to new information with minimal delay.

## 3.4 The Router Agent: Tool Selection without Decoding

The final component of the system is the Router Agent, which is responsible for selecting the next tool to execute based on the state of the "Mind." This is a simple, deterministic process that does not require any decoding or generation of text.

### 3.4.1 Reading the Intent Facet

The Router Agent's primary task is to read the state of the "intent" facet within the "Mind." This facet is a 512-dimensional vector that represents the system's current understanding of the user's goal. The Router Agent simply reads this vector from the SharedArrayBuffer, a fast and efficient operation.

### 3.4.2 Selecting the Next Tool Based on Vector State

Once the Router Agent has the "intent" vector, it can select the next tool to execute. This selection process can be implemented in several ways. The simplest approach is to compare the "intent" vector to a set of predefined tool vectors and select the tool with the highest cosine similarity.

### 3.4.3 Eliminating the Need for Prompts and Context Windows

The Router Agent's design completely eliminates the need for prompts and context windows. There is no need to construct a complex prompt to describe the user's intent to a generative model. The "intent" is already encoded in the "Mind" vector. This not only simplifies the system but also improves its performance, as it avoids the latency and computational overhead associated with prompt engineering and context management. The system operates in a purely vector-based, semantic space, making it highly efficient and responsive.

## 4. Experiments and Results

The performance of Neural Pin is evaluated based on several key metrics that highlight its efficiency and real-time capabilities. The results are compared against the operational characteristics of generative AI models to demonstrate the advantages of the post-generative paradigm.

### 4.1 Performance Metrics

The performance of Neural Pin is characterized by its speed, efficiency, and minimal resource consumption. The following metrics are used to quantify these characteristics.

#### 4.1.1 Convergence Time: 12 Human-Time Clicks ( $\approx 120$ ms)

The system's ability to converge on a user's intent is measured in terms of human-time clicks. The claim of convergence in 12 human-time clicks, which translates to approximately 120 milliseconds on a commodity laptop, is a powerful demonstration of its real-time performance. This metric is significant because it aligns the system's performance with the pace of human interaction, showing that it can learn and adapt at a speed that is imperceptible to the user.

#### 4.1.2 Energy Cost: $< 1$ mJ per Routing Decision

The energy efficiency of Neural Pin is a key design goal. The energy cost per routing decision is estimated to be less than 1 millijoule. This is orders of magnitude lower than the energy consumption of large generative models, which can require several kilojoules for a single inference pass. This low energy cost makes Neural Pin suitable for deployment on resource-constrained devices and in environments with strict power budgets.

#### 4.1.3 Code Size: 280 Lines of TypeScript + WebAssembly

The simplicity of the Neural Pin system is reflected in its small code size. The entire system is implemented in just 280 lines of TypeScript and WebAssembly. This minimalist implementation is a testament to the elegance of the post-generative paradigm and demonstrates that complex, real-time AI systems can be built with a relatively small amount of code.

### 4.2 Comparison with Generative AI Models

The performance of Neural Pin is starkly contrasted with that of generative AI models, highlighting the fundamental differences between the two paradigms.

#### 4.2.1 Latency and Throughput vs. Token-Based Systems

The latency of Neural Pin is several orders of magnitude lower than that of token-based systems. While a generative model may take seconds to produce a response, Neural Pin can update its state and make a routing decision in microseconds. This is because it avoids the sequential, iterative process of token generation. In terms of throughput, Neural Pin can handle a high velocity of incoming tool results without being bottlenecked by a slow generation process.

#### 4.2.2 Memory Footprint: 64 kB vs. KV Caches

The memory footprint of Neural Pin is a radical departure from that of generative models. The fixed 64 kB size is a tiny fraction of the memory required for the KV caches of large language models, which can easily reach gigabytes or even terabytes. This makes Neural Pin highly scalable and deployable on a wide range of hardware, from high-performance servers to resource-constrained edge devices.

#### 4.2.3 Energy Efficiency: Sub-millijoule vs. Large-Scale Models

The energy efficiency of Neural Pin is another key differentiator. The sub-millijoule energy cost per routing decision is a stark contrast to the high energy consumption of large-scale models. This is a direct result of the system's minimalist design and its avoidance of the computationally intensive operations that are at the heart of the transformer architecture.

### 4.3 Real-Time Responsiveness

Feature Neural Pin (Post-Generative) Generative AI (e.g., GPT-4)

Core Operation Stateful memory write Autoregressive token generation

Latency Sub-millisecond Hundreds of milliseconds to seconds

Memory Footprint 64 kB (fixed) Gigabytes to terabytes (KV cache)

Energy Cost < 1 mJ per decision Kilojoules per query

Learning Online, gradient-free update Offline, gradient-based training

Output Vector state (intent) Token sequence (text)

The real-time responsiveness of Neural Pin is a key feature that enables a new class of interactive AI applications.

#### 4.3.1 Human Reaction Time as a Benchmark

The use of human reaction time as a benchmark for convergence is a clever way to frame the system's performance. By showing that it can converge in less time than it takes for a human to perceive a response and initiate a new action, Neural Pin demonstrates that it can operate at the speed of human thought, enabling a fluid and natural interaction.

#### 4.3.2 Wall-Clock Time on a Commodity Laptop

The performance metrics of Neural Pin are measured on a commodity laptop, which is a realistic and relatable hardware environment. This demonstrates that the system does not require specialized or expensive hardware to achieve its impressive performance, making it accessible to a wide range of developers and users.

#### 4.3.3 Zero Allocations After Cold Start

The zero-allocation policy of Neural Pin is a critical factor in its real-time responsiveness. By avoiding dynamic memory allocation and garbage collection after the initial setup, the system ensures that its performance is consistent and predictable, without the latency spikes that can be caused by memory management in other systems.

## 5. Related Work

The design of Neural Pin is situated at the intersection of several distinct fields, including semantic search, real-time systems, and novel neural network architectures. Its core principles—a token-free, lock-free, and interrupt-driven approach to semantic routing—draw inspiration from and stand in contrast to existing methodologies in each of these areas.

### 5.1 Alternatives to Token-Based Generation

The generative AI landscape is overwhelmingly dominated by token-based models, particularly the Transformer architecture and its derivatives. However, the limitations of this approach have spurred research into alternative paradigms. Neural Pin represents a radical departure from this token-centric worldview, proposing a "post-generative" model where computation is driven by continuous vector states rather than discrete symbols.

#### 5.1.1 Retrieval-Augmented Generation (RAG) Models

Retrieval-Augmented Generation (RAG) has emerged as a powerful paradigm for enhancing the capabilities of large language models by grounding their responses in external, verifiable knowledge. A typical RAG system operates in a two-stage process:

first, a retriever component searches a large corpus of documents to find the most relevant information for a given user query; second, a generator component (usually a large language model) uses this retrieved information to construct a final, contextually informed response. The core of the retrieval stage often relies on semantic search, where both the user query and the documents in the corpus are encoded as high-dimensional vectors. The most relevant documents are then identified by computing the cosine similarity between the query vector and the document vectors, with the highest-scoring documents being selected for the generation stage. However, while RAG systems leverage vector similarity for retrieval, they remain fundamentally token-based in their operation. The generator component still processes the retrieved text as a sequence of tokens, and the overall system is still subject to the computational and memory overhead associated with large language models. Neural Pin, in contrast, eliminates the token-based generator entirely. It uses vector similarity not as a retrieval mechanism for a separate generative model, but as the primary computational primitive for driving the entire system.

#### 5.1.2 Sentence-BERT and Semantic Search

The use of pre-trained models like Sentence-BERT (SBERT) has become a standard practice in semantic search and RAG systems. SBERT is a modification of the original BERT model that adds a pooling layer to transform word-level embeddings into a single, fixed-size sentence-level embedding vector. This makes it particularly well-suited for tasks like semantic search and text similarity comparison, as the cosine similarity between these sentence vectors can be used as a direct measure of their semantic relatedness. While SBERT and similar models provide a powerful way to encode semantic information into vectors, they are still part of the broader Transformer ecosystem and inherit many of its characteristics. Neural Pin's "Mind" can be thought of as a highly specialized, continuously updated state vector that serves a similar purpose to a sentence embedding, but it is updated in real-time through a lock-free, online learning process, rather than being computed from a static input by a large, pre-trained model.

#### 5.1.3 Post-Generative Models and Non-Transformer Architectures

The term "post-generative" in the context of Neural Pin signifies a move away from the generative paradigm that defines most of modern AI. Instead of focusing on generating new content (tokens, images, etc.), the system's goal is to perform a specific action (routing to a tool) based on a continuously evolving understanding of the current context. This is a subtle but important distinction. While generative models are designed to be creative and produce novel outputs, Neural Pin is designed to be reactive and stateful, updating its internal representation of the world in response to new information and making decisions based on that state. This approach is more akin to a control system or an embedded agent than a traditional AI model. There is a growing body of research exploring non-Transformer architectures and alternative computational models for AI. For example, some work has focused on developing neural networks that do not rely on back-propagation for training, a technique known as "NoProp". These models aim to overcome some of the limitations of back-propagation, such as its biological implausibility and its computational cost. While

Neural Pin does use a form of gradient update, it is a highly simplified, in-place update that does not require the complex, multi-layer back-propagation algorithm used in deep learning.

## 5.2 Lock-Free and Real-Time Systems

The design of Neural Pin is heavily influenced by principles from the fields of real-time systems and concurrent programming. The use of a lock-free ring buffer, a shared memory architecture, and an interrupt-driven operational model are all hallmarks of high-performance, low-latency systems.

### 5.2.1 Lock-Free Data Structures in Shared Memory

Lock-free data structures are a key component of many high-performance concurrent systems. They allow multiple threads to access and modify a shared data structure without the need for traditional locking mechanisms like mutexes or semaphores. This can lead to significant performance improvements, as it eliminates the overhead of lock contention and the risk of deadlocks. The Spindle component of Neural Pin is a prime example of a lock-free data structure. It uses an atomic compare-and-swap (CAS) operation to update the head pointer of the ring buffer, ensuring that only one thread can write to a given slot at a time, without ever blocking another thread. This approach is particularly well-suited for the multi-tenant architecture of Neural Pin, where multiple "Minds" (one for each user or session) are all reading from the same shared ring buffer. The lock-free design ensures that the activity of one user cannot negatively impact the performance of another, which is a critical requirement for a responsive, interactive system.

### 5.2.2 Real-Time Operating Systems and Interrupts

The operational flow of Neural Pin is modeled after the interrupt-driven architecture of a real-time operating system (RTOS). In an RTOS, the main program is often idle, waiting for an interrupt to occur. When an interrupt is triggered by a hardware device (e.g., a timer, a sensor, or a network interface), the processor suspends its current task, saves its context, and jumps to a pre-defined interrupt service routine (ISR) to handle the event. Once the ISR is complete, the processor restores its context and resumes its previous task. The "Pin" in Neural Pin plays the role of the interrupt, and the "Mind" is the ISR. The system is idle, waiting for a new tool result to be written to the ring buffer. When the Spindle updates the head pointer, it triggers the Pin, which in turn wakes up the Mind. The Mind then performs its in-place gradient update, which is a very short, bounded-time operation, and then goes back to sleep. This interrupt-driven model is highly efficient, as it avoids the need for the Mind to constantly poll the ring buffer for new data, which would be a waste of CPU cycles.

### 5.2.3 Energy-Efficient Computing and Green AI

The energy efficiency of Neural Pin is one of its key design goals. The system is designed to be as lightweight as possible, with a fixed memory footprint of only 64 kB and a code size of just 280 lines of TypeScript and WebAssembly. The use of SIMD instructions for the vector updates and the lock-free, interrupt-driven architecture all contribute to its low energy consumption. The system is designed to do the minimum amount of work necessary to achieve its goal, and to be in a low-power idle state as much as possible. This focus on energy efficiency is part of a broader trend in the AI

community towards "Green AI," which aims to develop more sustainable and environmentally friendly AI systems. The massive computational and energy requirements of training and running large-scale generative models have raised concerns about their environmental impact. By demonstrating that a highly capable and responsive AI system can be built with a tiny fraction of the resources required by a large language model, Neural Pin provides a compelling example of what is possible with a more thoughtful and efficient approach to AI system design.

## 6. Conclusion

Neural Pin presents a compelling alternative to the dominant generative AI paradigm, demonstrating that a highly responsive and efficient AI system can be built on a foundation of simplicity and statefulness. By moving beyond token generation and treating inference as a stateful memory write, Neural Pin sidesteps the fundamental bottlenecks of latency, memory, and energy that plague large language models. The system's minimalist design, with its 64 kB memory footprint and lock-free architecture, makes it a practical solution for real-time, interactive applications.

### 6.1 Summary of Neural Pin's Contributions

The primary contributions of Neural Pin are threefold. First, it introduces a novel architectural pattern for AI systems that is fundamentally different from the dominant generative model. Second, it provides a practical, lightweight implementation of this architecture that can be deployed on commodity hardware. Third, it demonstrates the viability of this new paradigm through a series of experiments that showcase its real-time performance and low resource consumption. By challenging the conventional wisdom of token-based AI, Neural Pin opens up a new avenue of research and development, with the potential to enable a new generation of AI systems that are more agile, more efficient, and more deeply integrated into the real world.

### 6.2 Implications for Post-Generative AI

The success of Neural Pin has significant implications for the future of AI. It suggests that the path to more capable and efficient AI may not lie in simply building larger and more complex generative models, but in exploring alternative paradigms that are better suited to the demands of real-world applications. The post-generative approach, with its focus on statefulness, efficiency, and real-time responsiveness, offers a promising direction for future research. It points towards a future where AI is not just a tool for generating content, but a deeply integrated partner in our daily lives, capable of understanding our intent and acting on our behalf with speed and precision.

### 6.3 Future Work and Potential Applications

The work on Neural Pin opens up several avenues for future research and development. One potential area of exploration is the application of the post-generative paradigm to other domains, such as robotics, autonomous systems, and personalized healthcare. The system's low latency and energy efficiency make it particularly well-suited for these applications. Another area of future work is the development of more sophisticated learning rules and state update mechanisms. While the current system uses a simple gradient descent rule, more advanced techniques could be explored to improve the system's learning capabilities and adaptability. Finally, the integration of Neural Pin with generative models could lead to hybrid systems that combine the real-



time responsiveness of the post-generative approach with the creative power of generative AI. Such systems could offer the best of both worlds, enabling a new class of AI applications that are both highly interactive and deeply creative.

Bryan Conklin 10/11/2025